



HPC application design: ParFlow GPU support

2022-07-05 | Jaro Hokkanen



What is ParFlow?

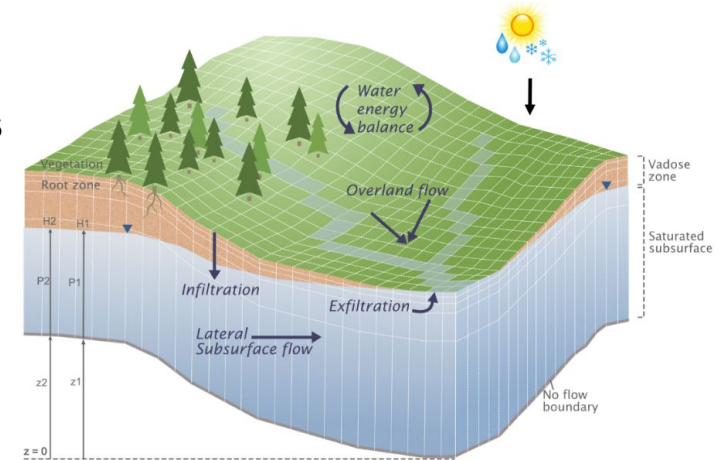
- 3D model for variably saturated subsurface flow including pumping and irrigation, and integrated overland flow
(Jones & Woodward, 2001; Kollet & Maxwell, 2006; Maxwell, 2013)
- Integrated with land surface and also regional climate model, TSMP (Shrestha et al., 2014)
- External coupling via OASIS3-MCT (Shrestha et al., 2014; Gasper et al., 2014)
- Parallel Data Assimilation Framework with PDAF (Kurtz et al., 2016)
- Optimized for modern heterogeneous supercomputers environments
(e.g. Gasper et al., 2014, Burstedde et al., 2018; Hokkanen et al., 2021)



https://github.com/hydroframe/ParFlow_Short_Course/blob/master/Slides/1.parflow_intro.pdf

ParFlow numerical model

- Cell centered finite difference / finite volume schemes
- Implicit time integration
- Newton-Krylov methods (Kinsol) with multigrid preconditioning (ParFlow, Hypre) for nonlinear problems
- A large number of numerical kernels, none of which clearly dominate the run time
- Around 150k lines of C



https://github.com/hydroframe/ParFlow_Advanced_ShortCourse/blob/master/slides/1.Intro_Overview.pdf



Desirable requirements for the accelerator support

- Good performance
- Single-source application
- Codebase long-term maintainability
- Minimum reliance on external dependencies
- Separation of concerns (scientific development vs. accelerator utilization)
- High developer productivity for adding support for new accelerators

- ParFlow solution: Support for accelerators is added into a macro-based abstraction layer (ParFlow embedded Domain-Specific Language, ParFlow eDSL)

ParFlow eDSL interface

Allocations & initializations

```
KW = NewVectorType(grid2d, 1, 1, cell_centered);  
InitVector(KW, 0.0);
```

Message passing

```
update_handle = InitVectorUpdate(  
    pressure, VectorUpdateAll);  
FinalizeVectorUpdate(update_handle);
```

Accessor macros

```
ix = SubgridIX(subgrid);  
iy = SubgridIY(subgrid);  
iz = SubgridIZ(subgrid);
```

Loop macros

```
GrGeomInLoop(i, j, k, gr, r, ix, iy, iz, nx, ny, nz,  
{  
    int ips = SubvectorEltIndex(ps_sub, i, j, k);  
    data[ips] = value;  
});
```

ParFlow eDSL for GPUs

- ParFlow can use a native CUDA implementation or the Kokkos library for GPUs
 - Both use the same interface for memory management and compute kernels
- Some of the most recent technology is used for best performance and developer productivity:
 - Unified Memory
 - Host-device lambdas
 - Pool allocator for Unified Memory
- CUDA-aware MPI library and GPU-based application-side data packing routines can be used for fast GPU-GPU communication

ParFlow eDSL for GPUs: Incremental development

- In ParFlow, memory (de)allocations/compute kernels are accessed through the eDSL
- Unified Memory (de)allocations/compute kernels are controlled for each compilation unit separately allowing incremental development and flexibility
 - Some source files use Unified Memory and GPU-based loops, and others not

```
/* PFCUDA_COMP_UNIT_TYPE determines the compilation unit type:  
  1: HIP/NVCC compiler, Unified Memory allocation, Parallel loops (GPUs)  
  2: HIP/NVCC compiler, Unified Memory allocation, Sequential loops (CPUs)  
  default: HIP/NVCC compiler, Standard heap allocation, Sequential loops (CPUs) */  
#define PFCUDA_COMP_UNIT_TYPE 1 // Defined by CMake
```

Memory management: Host memory

Source file

```
vector = talloc(Vector, 1);
```

```
tfree(vector);
```

eDSL header file

```
#define talloc(type, count) \  
    (type*)malloc(sizeof(type) \  
        * (unsigned int)(count))
```

```
#define tfree(ptr) free(ptr)
```


Memory management: Unified Memory

CPU version eDSL header file

```
#define talloc(type, count) \  
    (type*)malloc(sizeof(type) \  
        * (unsigned int)(count))
```

GPU version eDSL header file

```
#define talloc(type, count) \  
    (type*)talloc_hip(sizeof(type) \  
        * (unsigned int)(count))  
  
static inline void *talloc_hip(size_t size)  
{  
    void *ptr = NULL;  
    // hipMallocManaged(&ptr, size);  
    rmmAlloc(&ptr, size, 0, __FILE__, __LINE__);  
    return ptr;  
}
```

Memory management: Unified Memory

CPU version eDSL header file

```
#define tfree(ptr) free(ptr)
```

GPU version eDSL header file

```
#define tfree(ptr) tfree_hip(ptr)

static inline void tfree_hip(void *ptr)
{
    // hipFree(&ptr);
    rmmFree(ptr, 0, __FILE__, __LINE__);
}
```

Loops: An example of a specialized GPU kernel

Original source file

```
BoxLoopI0(i, j, k, ix, iy, iz, nx, ny, nz,
{
    int ip;
    ip = SubvectorEltIndex(f_sub, i, j, k);
    fp[ip] = pp[ip] - value;
});
```

What could be done, but...

```
#ifndef HAVE_HIP
/* some code to find grid & block sizes */
HIPKernel<<<grid, block>>>(ix, ...);
#else
BoxLoopI0(i, j, k, ix, iy, iz, nx, ny, nz,
{
    int ip;
    ip = SubvectorEltIndex(f_sub, i, j, k);
    fp[ip] = pp[ip] - value;
});
#endif
```

Loops: General GPU kernels

Original source file

```
/* ... using CPU macros ... */  
BoxLoopI0(i, j, k, ix, iy, iz, nx, ny, nz,  
{  
    int ip;  
    ip = SubvectorEltIndex(f_sub, i, j, k);  
    fp[ip] = pp[ip] - value;  
});
```

New source file

```
/* ... using GPU macros ... */  
BoxLoopI0(i, j, k, ix, iy, iz, nx, ny, nz,  
{  
    int ip;  
    ip = SubvectorEltIndex(f_sub, i, j, k);  
    fp[ip] = pp[ip] - value;  
});
```

Loops: General GPU kernels

CPU version eDSL header file

```
#define BoxLoopI0(i, j, k, \
    ix, iy, iz, nx, ny, nz, loop_body) \
{ \
    for (k = iz; k < iz + nz; k++) \
        for (j = iy; j < iy + ny; j++) \
            for (i = ix; i < ix + nx; i++) \
                { \
                    loop_body; \
                } \
}
```

GPU version eDSL header file (CUDA/HIP)

```
#define BoxLoopI0(i, j, k, \
    ix, iy, iz, nx, ny, nz, loop_body) \
{ \
    /* some code to find grid & block sizes */ \
    auto lambda_body = [=] __host__ __device__ \
        (int i, int j, int k) \
        { \
            i += ix; j += iy; k += iz; \
            loop_body; \
        } \
    BoxKernelI0<<<grid, block>>>(lambda_body, \
        nx, ny, nz); \
}
```

Loops: General GPU kernels

GPU version eDSL header file (CUDA/HIP) GPU version eDSL header file (Kokkos)

```
#define BoxLoopI0(i, j, k, \
    ix, iy, iz, nx, ny, nz, loop_body) \
{ \
    /* some code to find grid & block sizes */ \
    auto lambda_body = [=] __host__ __device__ \
        (int i, int j, int k) \
    { \
        i += ix; j += iy; k += iz; \
        loop_body; \
    } \
    BoxKernelI0<<<grid, block>>>(lambda_body, \
        nx, ny, nz); \
}
```

```
#define BoxLoopI0(i, j, k, \
    ix, iy, iz, nx, ny, nz, loop_body) \
{ \
    auto lambda_body = KOKKOS_LAMBDA \
        (int i, int j, int k) \
    { \
        i += ix; j += iy; k += iz; \
        loop_body; \
    } \
    MDPolicyType_3D mdpolicy_3d({{0, 0, 0}}, \
        {{nx, ny, nz}}); \
    Kokkos::parallel_for(mdpolicy_3d, \
        lambda_body); \
}
```

Loops: A general HIP/CUDA kernel example

```
template <typename LambdaBody>
__global__ static void BoxKernelI0(LambdaBody loop_body,
    const int nx, const int ny, const int nz)
{
    int i = ((blockIdx.x * blockDim.x) + threadIdx.x);
    int j = ((blockIdx.y * blockDim.y) + threadIdx.y);
    int k = ((blockIdx.z * blockDim.z) + threadIdx.z);

    if(i < nx && j < ny && k < nz)
    {
        loop_body(i, j, k);
    }
}
```

A complete hello world example



```
#include <stdio.h>

/* Kernel macro selection (CPU/GPU) */
// #define BoxLoop BoxLoopCPU
#define BoxLoop BoxLoopGPU

/* Compute kernel macro for CPU (API) */
#define BoxLoopCPU(i, nx, loop_body) \
{ \
    for (i = 0; i < nx; i++) \
    { \
        loop_body; \
    } \
}

/* Compute kernel macro for GPU (API) */
#define BoxLoopGPU(i, nx, loop_body) \
{ \
    auto lambda = [=] __host__ __device__ (int i) \
    { \
        loop_body; \
    }; \
 \
    int blocksize = 1024; \
    int gridsize = (nx - 1 + blocksize) / blocksize; \
 \
    _BoxKernel<<<gridsize, blocksize>>>(lambda, nx); \
    cudaStreamSynchronize(0); \
    (void)i; \
} \
}
```

```
/* General GPU kernel */
template <typename LambdaBody> __global__ static
void _BoxKernel(LambdaBody lambda, const int nx)
{
    const int i =
        blockIdx.x * blockDim.x + threadIdx.x;
    if(i < nx)
    {
        lambda(i);
    }
}

/* Memory allocation macro (API) */
#define alloc_managed(type, count) \
    (type*)_alloc_managed(count * sizeof(type));

/* Function to allocate Unified Memory */
static inline void *_alloc_managed(size_t size)
{
    void *ptr = NULL;
    hipMallocManaged(&ptr, size);
    return ptr;
}

/* Memory deallocation macro (API) */
#define free_managed(ptr) _free_managed(ptr);

/* Function to deallocate Unified Memory */
static inline void _free_managed(void *ptr)
{
    hipFree(ptr);
}
```

← Headers

Driver function



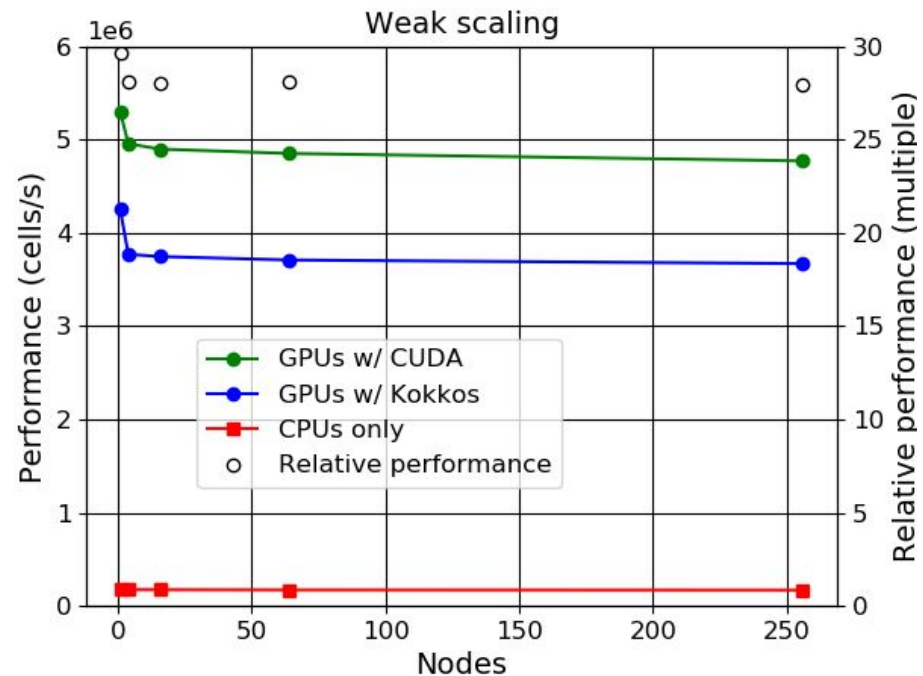

```
/* Driver function */
int main(int argc, char *argv [])
{
    int i, nx = 10;
    int* array = alloc_managed(int, nx);

    BoxLoop(i, nx,
    {
        array[i] = i;
    });

    BoxLoop(i, nx,
    {
        int thread = array[i];
        printf("Hello from GPU thread %d\n", thread);
    });

    free_managed(array);
}
```


Results from JUWELS Booster supercomputer



Summary

- The ParFlow results from JUWELS Booster supercomputer demonstrate a very good weak scaling with up to 28X speedup from A100 GPUs over hundreds of nodes
- Hiding GPU support into a macro-based DSL resulted in good developer productivity, codebase long-term maintainability, and performance gain
- An existing DSL is not a prerequisite: Only a minimal abstraction layer for accessing the most relevant loops and memory (de)allocations is required
- Unified Memory support played an important role in good developer productivity
- Pool allocator for GPU memory can play a very important role for high performance gain
- Does the approach work for Fortran? -Yes, but only with CUDA: The approach has been tested in COSMO atmospheric model using CUF kernels instead of lambda functions